

BETA

# Hardware- Based Security

Alessandro Savino

---



Politecnico  
di Torino

# Principles of Computer Architecture

Traditional computer architecture has six principles regarding processor design:

- Caching

E.g. caching frequently used data in a small but fast memory helps hide data access latencies.

- Pipelining

E.g. break processing of an instruction into smaller chunks that can each be executed sequentially reduces critical path of logic and improves performance.

- Predicting

E.g. predict control flow direction or data values before they are actually computed allows code to execute speculatively.

- Parallelizing

E.g. processing multiple data in parallel allows for more computation to be done concurrently.

- Use of indirection

E.g. virtual to physical mapping abstracts away physical details of the system.

- Specialization

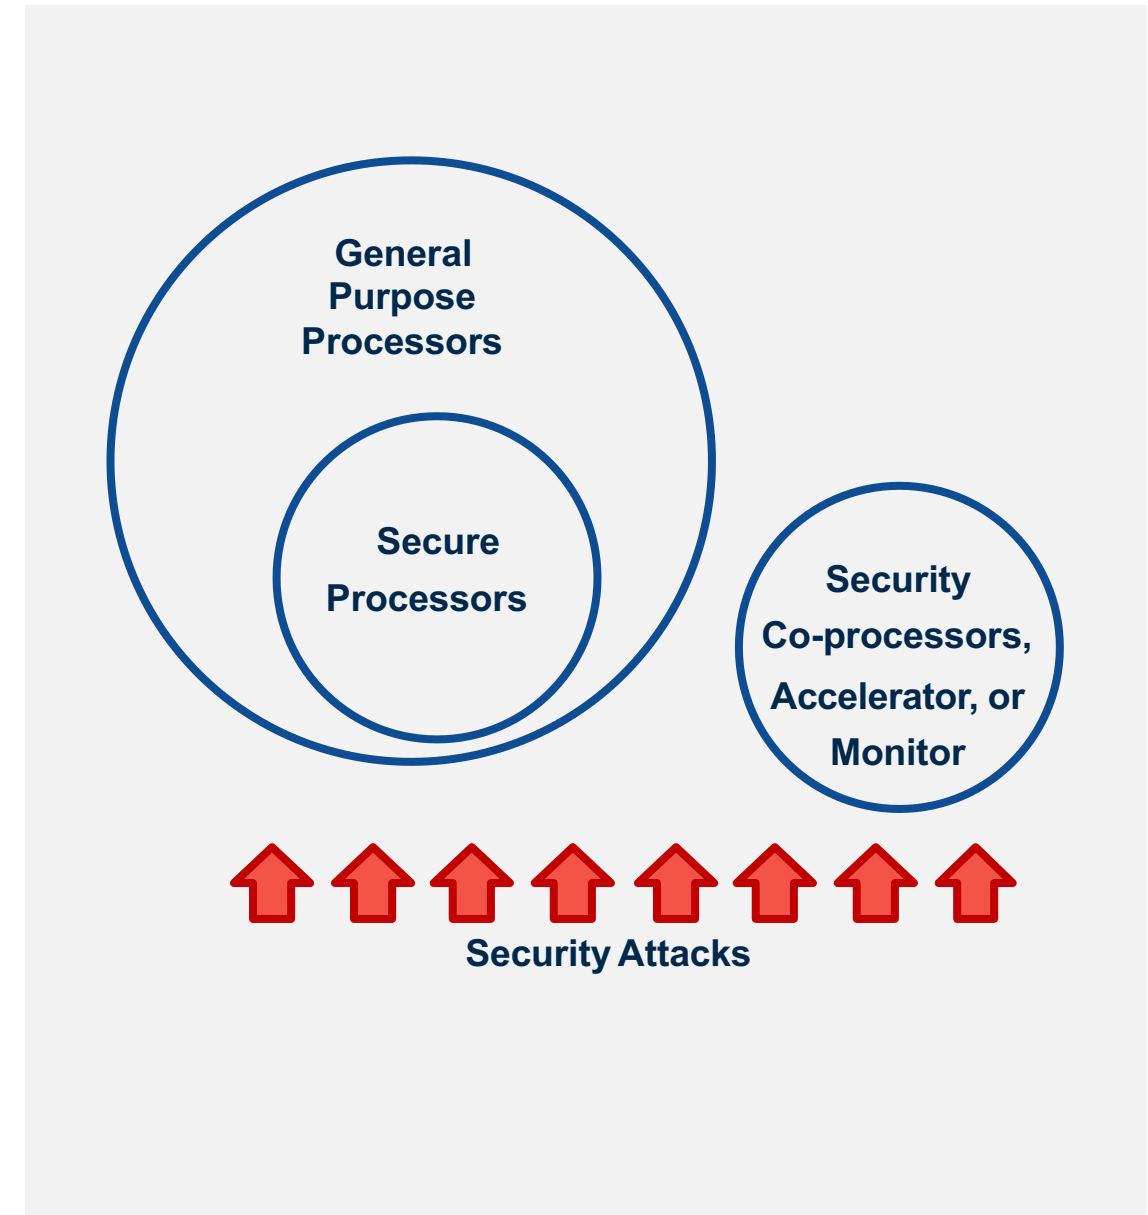
E.g. custom instructions use dedicated circuits to implement operations that otherwise would be slower using regular processor instructions.

**What are principles for securing processors?**

# Processor Security and Secure Processors

Secure processors:

- › Subset of processors with extra security features
- › Provide extra logical isolation for software
- › Vulnerable to similar attacks as regular processors



# Secure Processor Architectures

Secure Processor Architectures extend a processor with hardware (and related software) features for the protection of software

Protected pieces of code and data are now commonly called Enclaves

- › But can also be Trusted Software Modules, whole Operating Systems, or Virtual Machines

Focus on the main processor in the system

- › Others focus on co-processors, cryptographic accelerators, or security monitors

Add more features to isolate secure software from other untrusted software

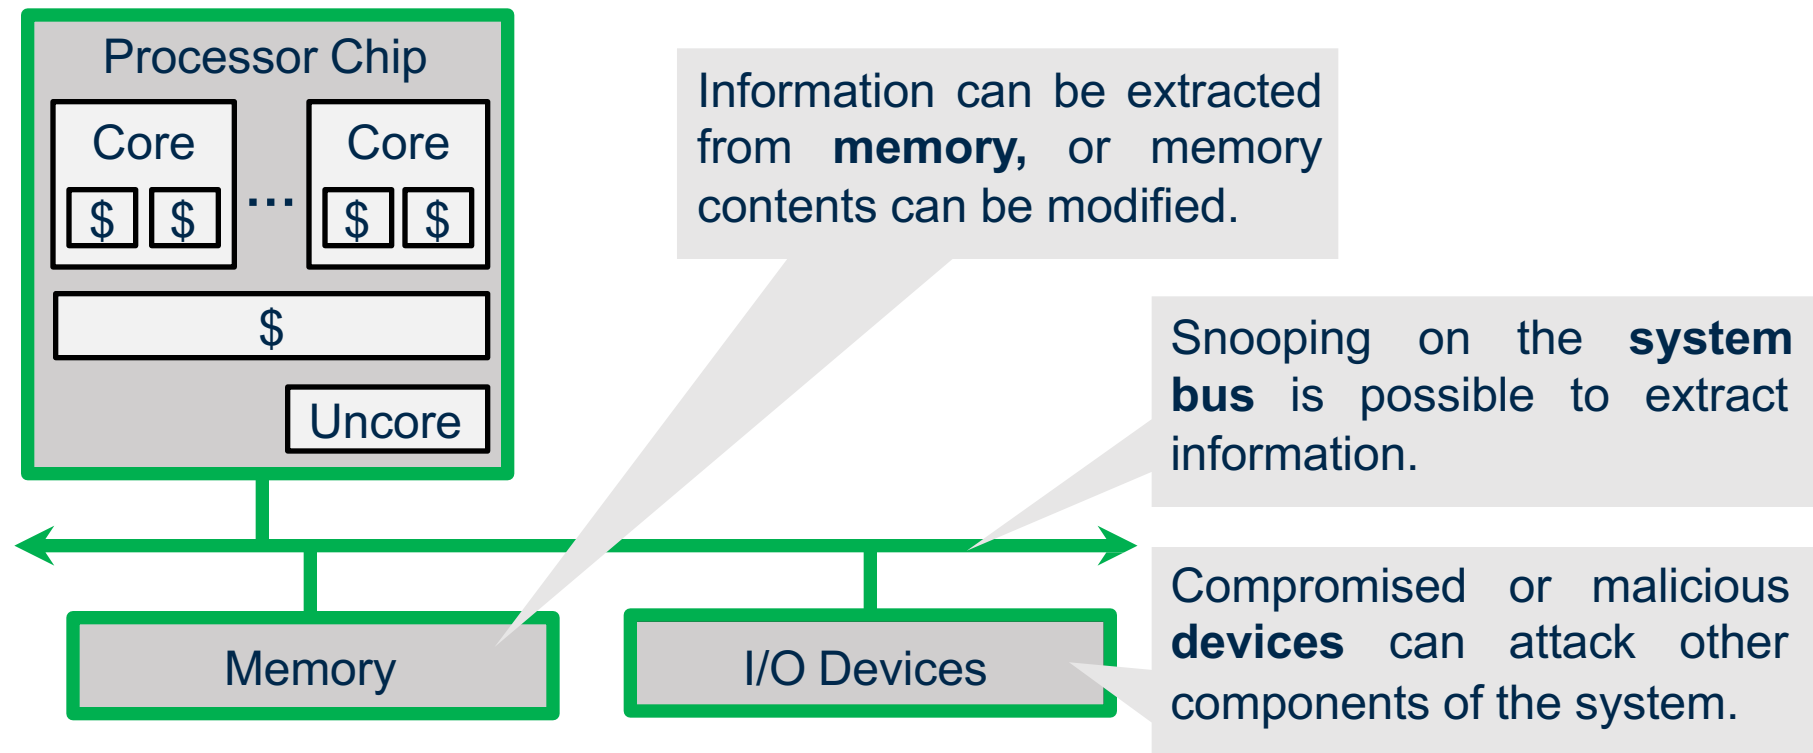
- › Includes untrusted Operating System or Virtual Machines
- › Many also consider physical attacks on memory

Isolation should cover all types of possible ways for information leaks

- › Architectural state
- › Micro-architectural state 
- › Due to spatial or temporal sharing of hardware 

# Baseline (Unsecure) Processor Hardware

A typical computer system with no secure components nor secure processor architectures considers all the components as trusted:



# Baseline (Unsecure) Processor Software

The typical computer system uses a ring-based protection scheme, which gives the most privileges (and most trust) to the lowest levels of the system software:




Image:  
[https://commons.wikimedia.org/wiki/File:Priv\\_rings.svg](https://commons.wikimedia.org/wiki/File:Priv_rings.svg)

# Logical Isolation with New Privilege Levels

Modern computer systems define protections regarding privilege levels or protection rings, and new privilege levels are defined to provide added protections.

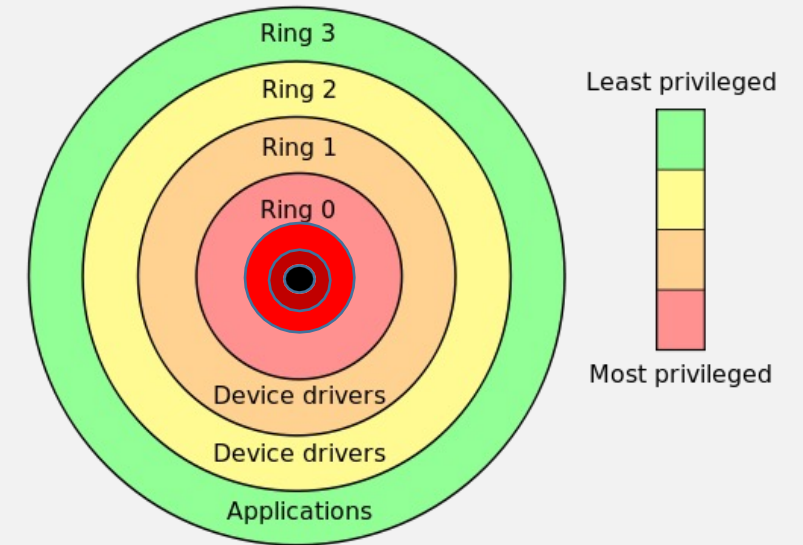


Image:

[https://commons.wikimedia.org/wiki/File:Priv\\_rings.svg](https://commons.wikimedia.org/wiki/File:Priv_rings.svg)

# Logical Isolation with New Privilege Levels

Ring 3: Application code, least privileged.

Ring 2 and 1: Device drivers and other semi-privileged code, although rarely used.

Ring 0: Operating system kernel.

Ring -1: Hypervisor or virtual machine monitor (VMM), the most privileged mode a typical system administrator can access.

Ring -2: System management mode (SMM), typically locked down by the processor manufacturer

Ring -3: The platform management engine, retroactively named “ring -3”, runs on a separate management processor.




Image:

[https://commons.wikimedia.org/wiki/File:Priv\\_rings.svg](https://commons.wikimedia.org/wiki/File:Priv_rings.svg)




# Extend Linear Trust to the New Protection Levels

The hardware is most privileged as it is the lowest level in the system.

- › There is a linear relationship between the protection ring and privilege (the lower ring is more privileged)
- › Each component trusts all the software “below” it

**Security Engine (SecE)** can be something like Intel’s ME or AMD’s PSP.



# Add Horizontal Privilege Separation

New privileges can be made orthogonal to existing protection rings.

- › E.g. , ARM's TrustZone's “normal” and “secure” worlds
- › Need privilege level (ring number) normal / secure privilege

Security levels from a lattice:






Image:

[https://commons.wikimedia.org/wiki/File:Priv\\_rings.svg](https://commons.wikimedia.org/wiki/File:Priv_rings.svg)

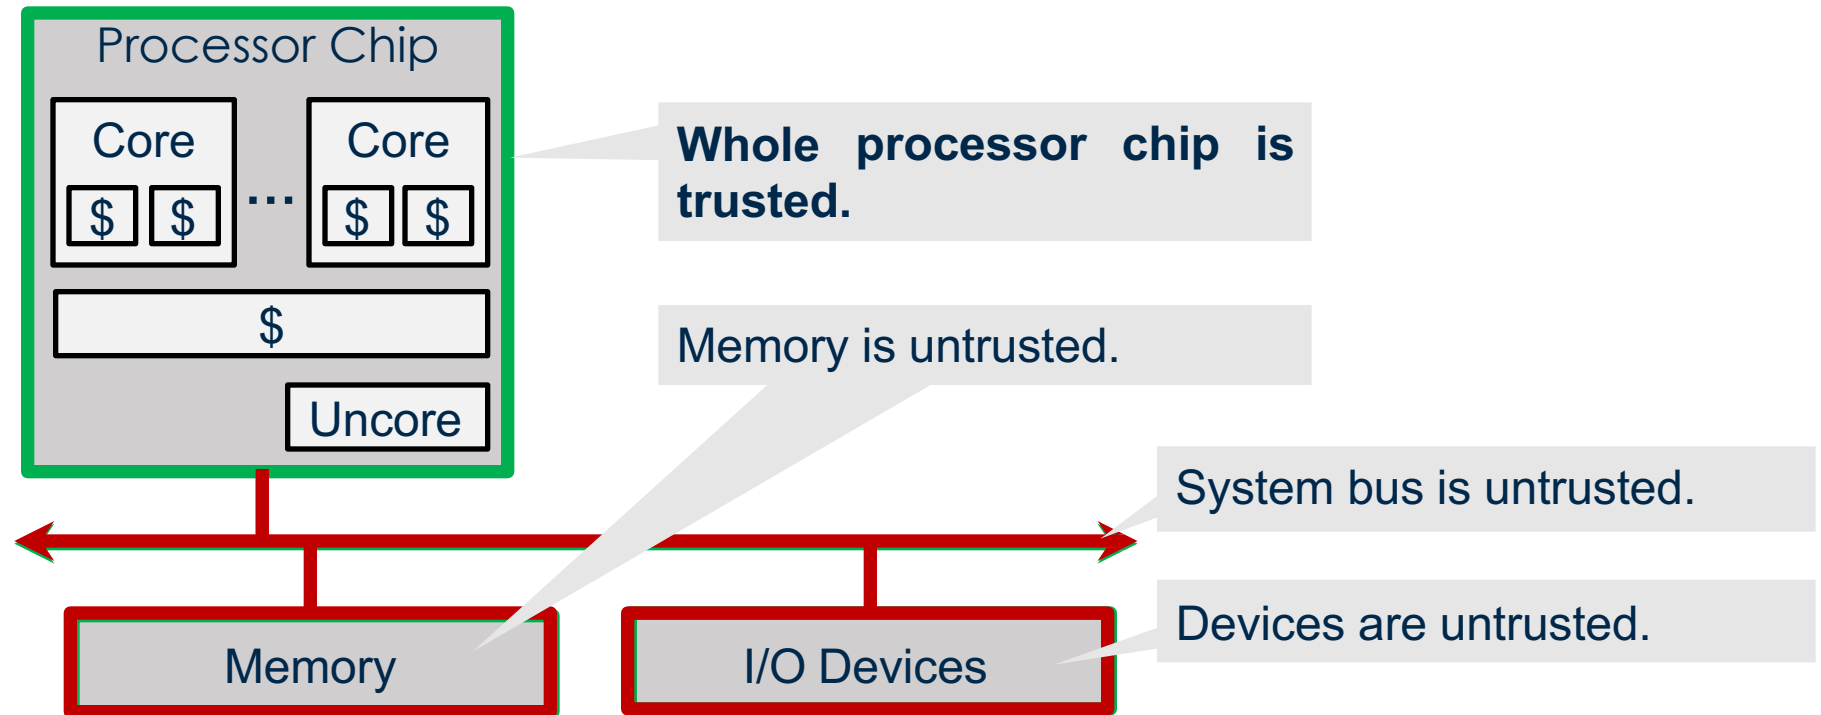
# Breaking Linear Hierarchy of Protection Rings

Examples of architectures that do and don't have a linear relationship between privileges and protection ring level:



# Providing Protections with a Trusted Processor Chip

Key to most secure processor architecture designs is the idea of a **trusted processor chip** as the security wherein the protections are provided.



# Limitations of Secure Processors

Threats that are outside the scope of secure processor architectures:

- › Bugs or Vulnerabilities in the TPC
- › Hardware Trojans and Supply Chain Attacks
- › Physical Probing and Invasive Attacks

Threats that are underestimated when designing secure processor architectures:

- › Side Channel Attacks
  - › Information can leak through timing, power, or electromagnetic emanations from the implementation

# TEE and TCB

The Trusted Computing Base (TCB) is the set of hardware and software that is responsible for realizing the Trusted Execution Environments (TEEs) :

- › TEE is created by a set of all the components in the TCB
- › TCB is trusted to implement the protections correctly
  - › TCB may not be trustworthy if it is not verified or is not bug-free
- › Vulnerability or successful attack on TCB nullifies TEE protections

The goal of TEE is to protect a piece of code and data from a range of software and hardware attacks.


Multiple mutually untrusting pieces of protected code can run on a system at the same time

# TEEs and Software They Protect

Different architectures mainly focus on protecting Trusted Software Modules (a.k.a. enclaves) or whole Virtual Machines.

Other TEEs support Trusted Software Modules, a.k.a. enclaves

Some TEEs have support for protecting whole virtual machines.



# Protections Offered by Secure Processor Architectures

Security properties for the TEEs that secure processor architectures aim to provide:

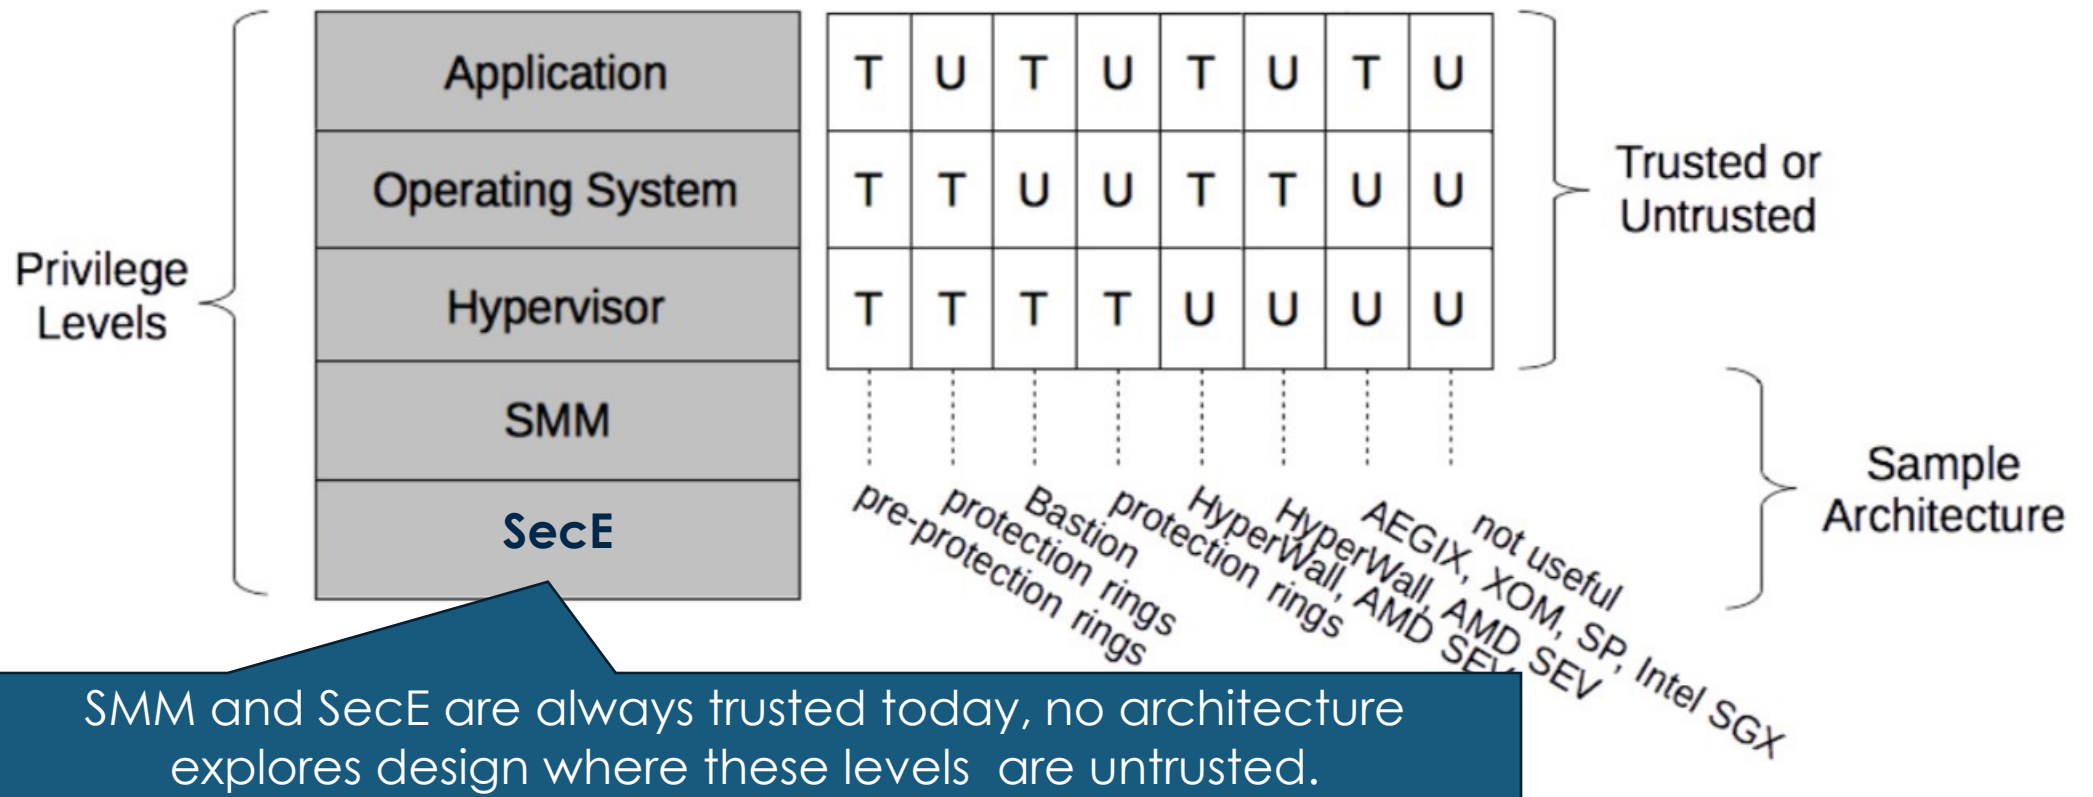
- › Confidentiality
- › Integrity
- › Availability is usually not provided usually

Confidentiality and integrity protections are from attacks by other components (and hardware) not in the TCB. There is typically no protection from malicious TCB.



# Protections Categorized by Architecture

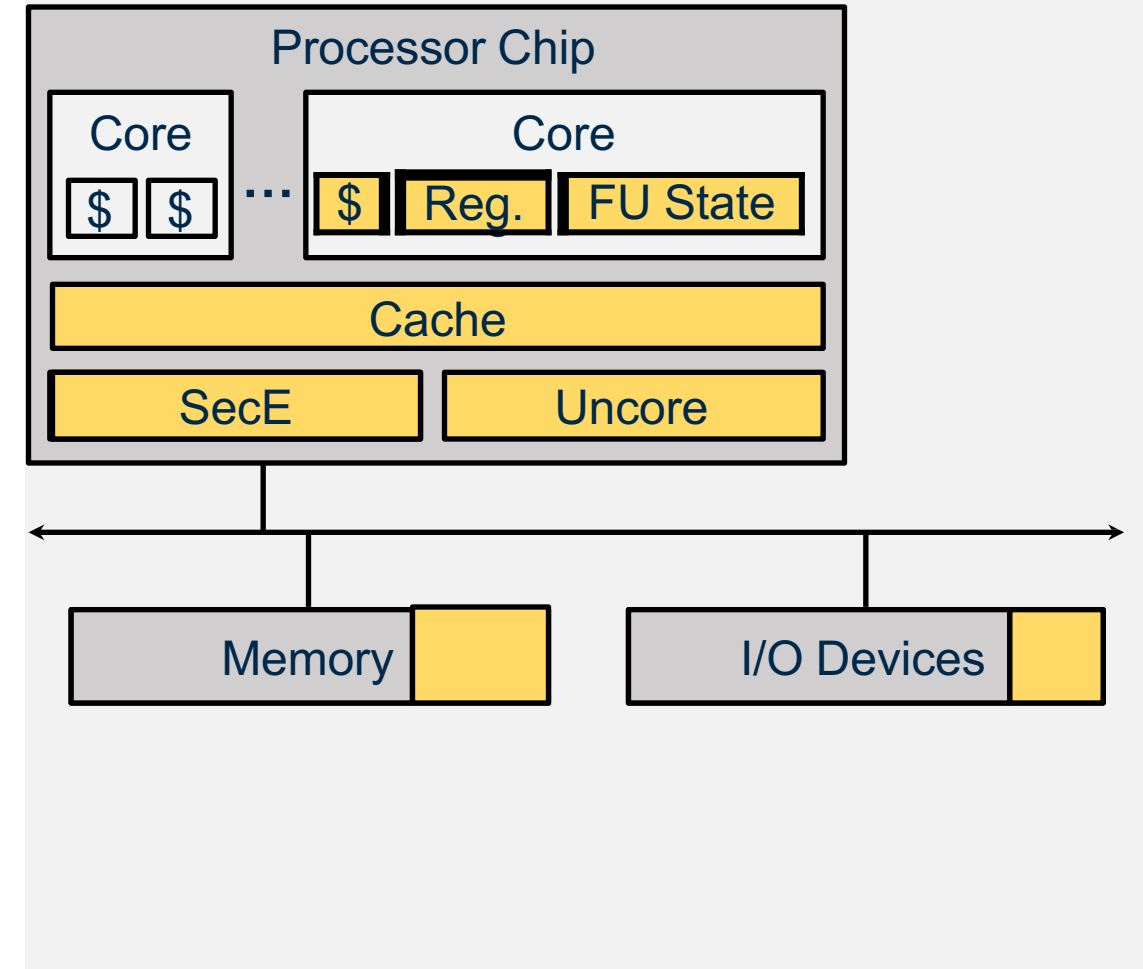
Secure processor architectures break the linear relationship (where lower level protection ring is more trusted):



# Protecting State of the Protected Software

The protected software's state is distributed throughout the processor. All of it needs to be protected from untrusted components and other (untrusted) protected software.

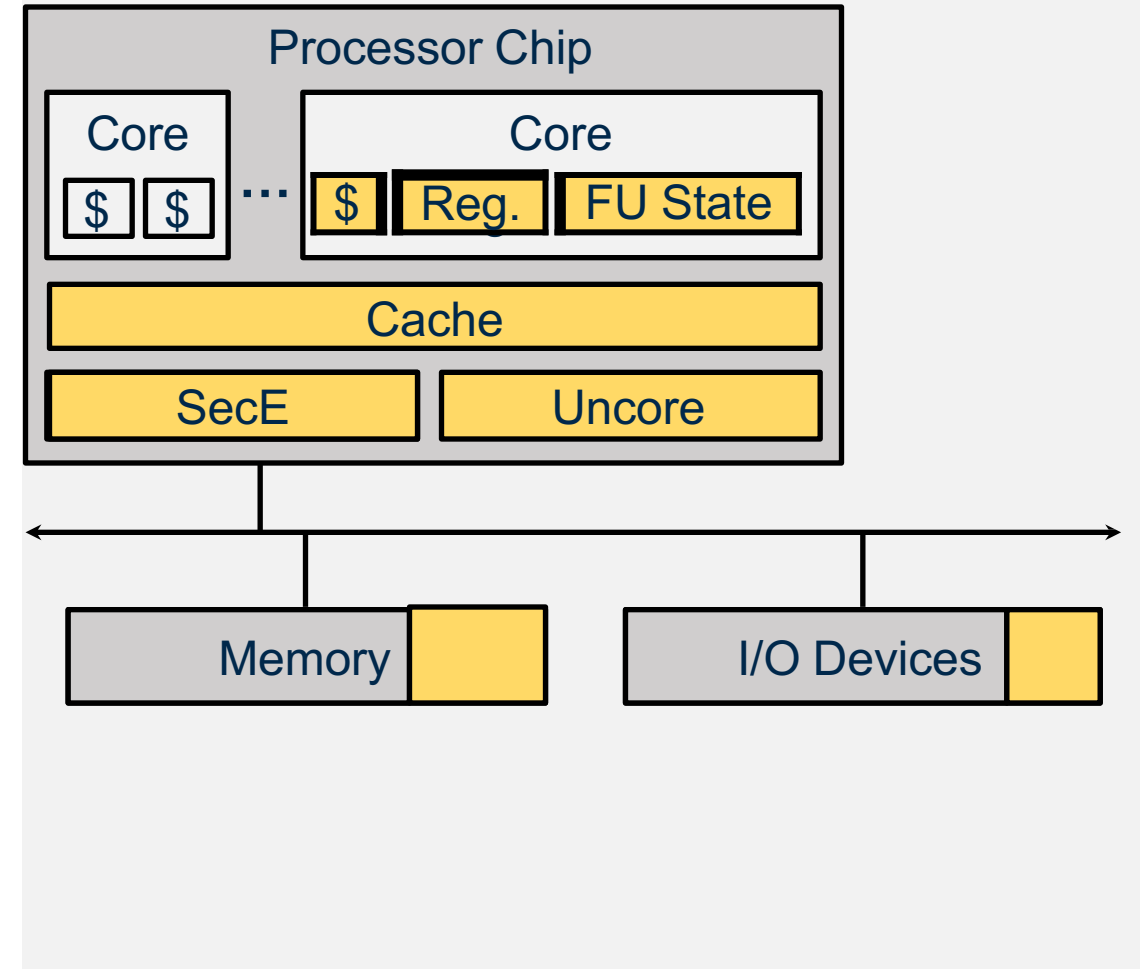
- › Protect memory through encryption and hashing with integrity trees
- › Flush state, or isolate state, of functional units inside processor cores
- › Isolate state in uncore and any security modules
- › Isolate state in I/O and other subsystems



# Ideal No Side-Effects Execution

Secure processor architectures ideally have no side effects, which are visible to the untrusted components whenever protected software is being executed.

- › The system is in some state before protected software runs
- › Protected software runs, modifying system state
- › When protected software is interrupted or terminates, the state modifications are erased



# No Protections from Protected Software

The software (code and data) executing within TEE protections is assumed to be benign and not malicious:

- › The goal of Secure Processor Architectures is to create minimal TCB that realizes a TEE within which the protected software resides and executes
- › Secure Processor Architectures cannot protect software if it is buggy or has vulnerabilities

Code bloat endangers invalidating assumptions about benign protected software.

Attacks from within protected software should be defended.

# Hardware TCB as Circuits or Processors

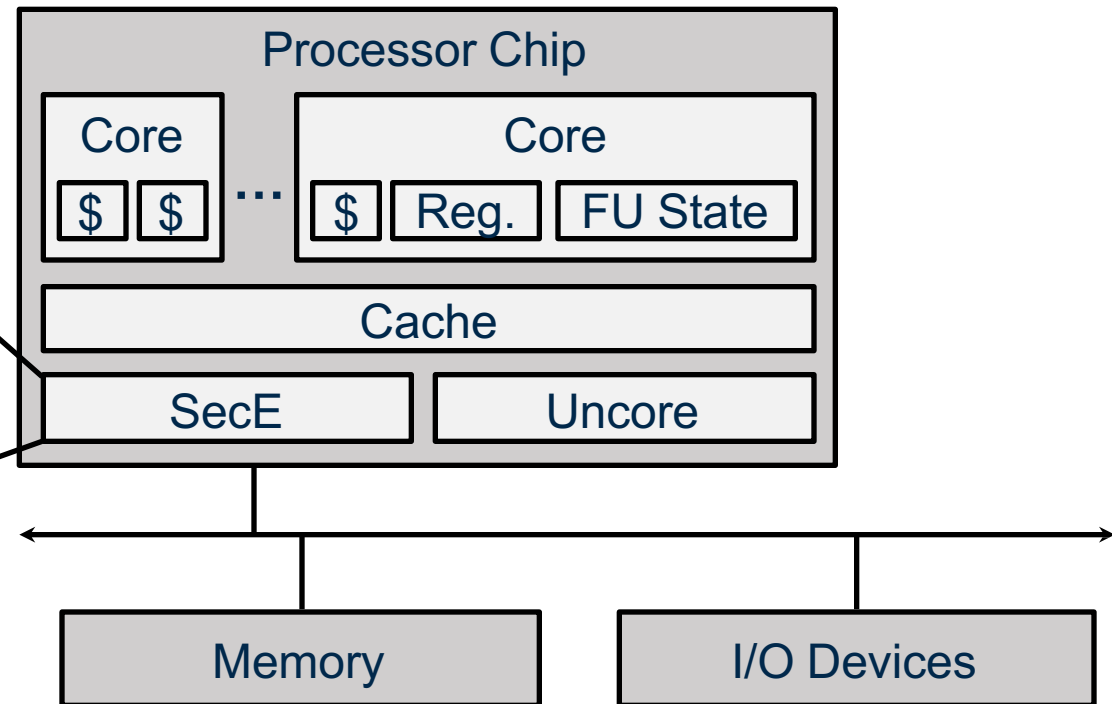
Key parts of the hardware TCB can be implemented as dedicated circuits or as firmware or other code running on dedicated processor

- **Custom logic or hardware state machine:**

- Most academic proposals

- **Code running on dedicated processor:**

- Intel ME = ARC processor or Intel Quark processor
- AMD PSP = ARM processor



**Vulnerabilities in TCB “hardware” can lead to attacks that nullify the security protections offered by the system.**

# Ensuring Trustworthy TCB Execution

The trustworthiness of the TCB depends on the ability to monitor the TCB code (hardware and software) execution as the system runs.

TCB should be monitored to ensure it is trustworthy.

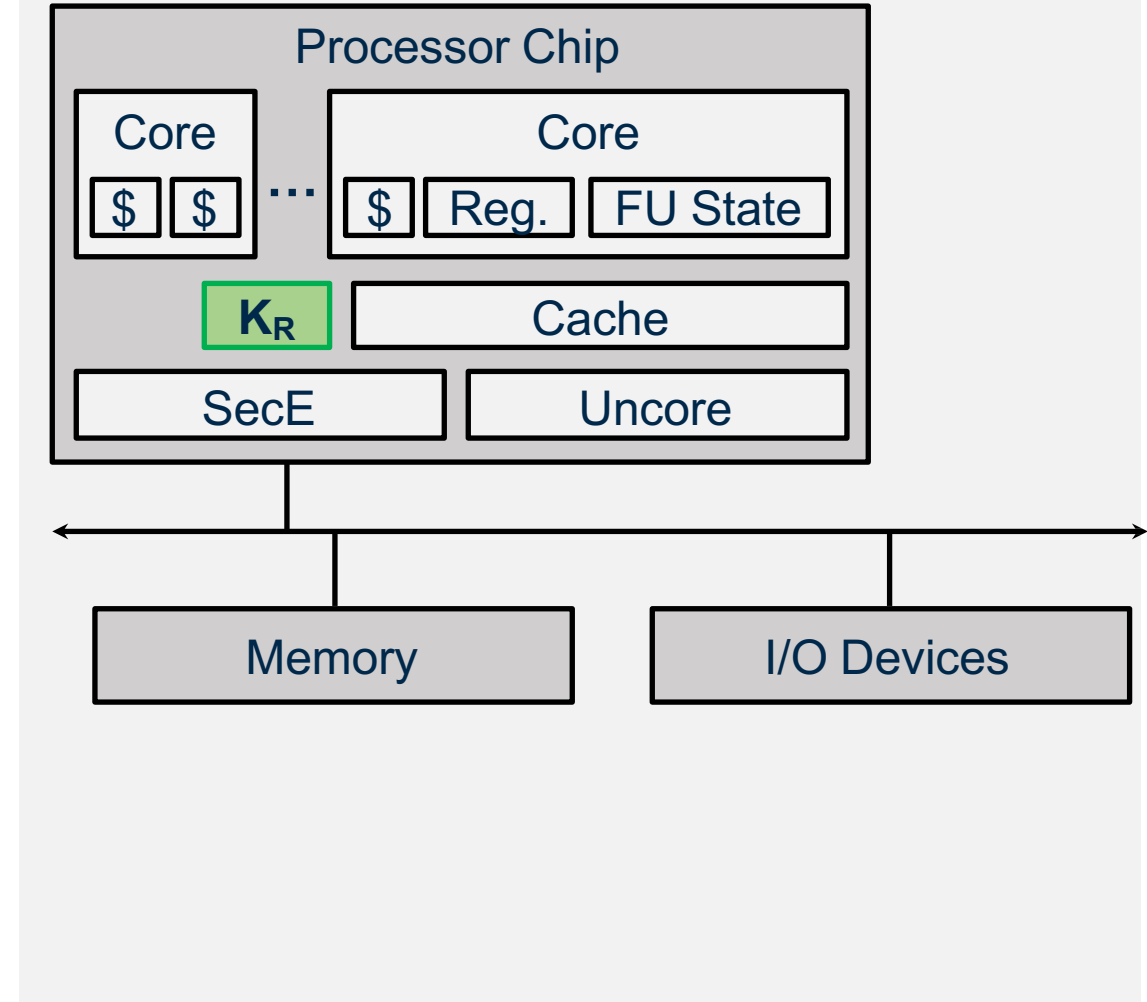
Monitoring of TCB requires mechanisms to:

- › Fingerprint and authenticate TCB code
- › Monitor TCB execution
- › Protect TCB code (on embedded security processor)
  - › Virtual Memory, ASLR, ...

# Root of Trust and Processor Key

Security of the system is derived from a root of trust.


- › A secret (cryptographic key) only accessible to TCB components
- › Derive encryption and signing keys from the root of trust



# Root of Trust for TCB

Each processor requires a unique secret:

- › Burn in at the factory by the manufacturer (but implies trust issues with the manufacturer and the supply chain)
  - › E.g., One-Time Programmable (OTP) fuses
- › Use Physically Unclonable Functions (but require reliability)
  - › Extra hardware to derive keys from PUF
  - › Mechanisms to generate and distribute certificates for the key

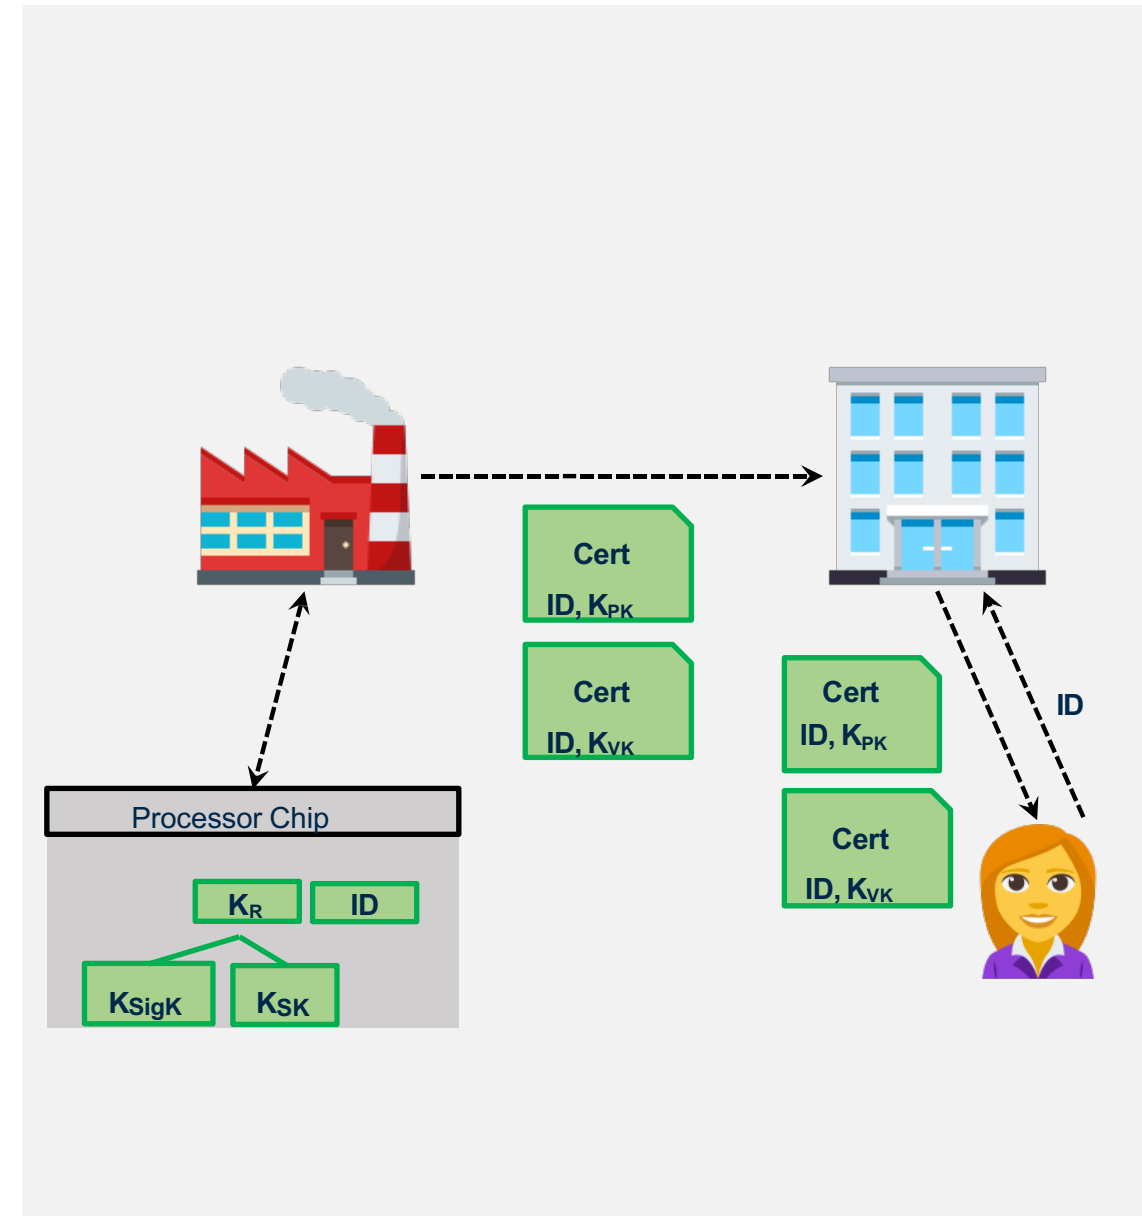




# Derived Keys and Key Distribution

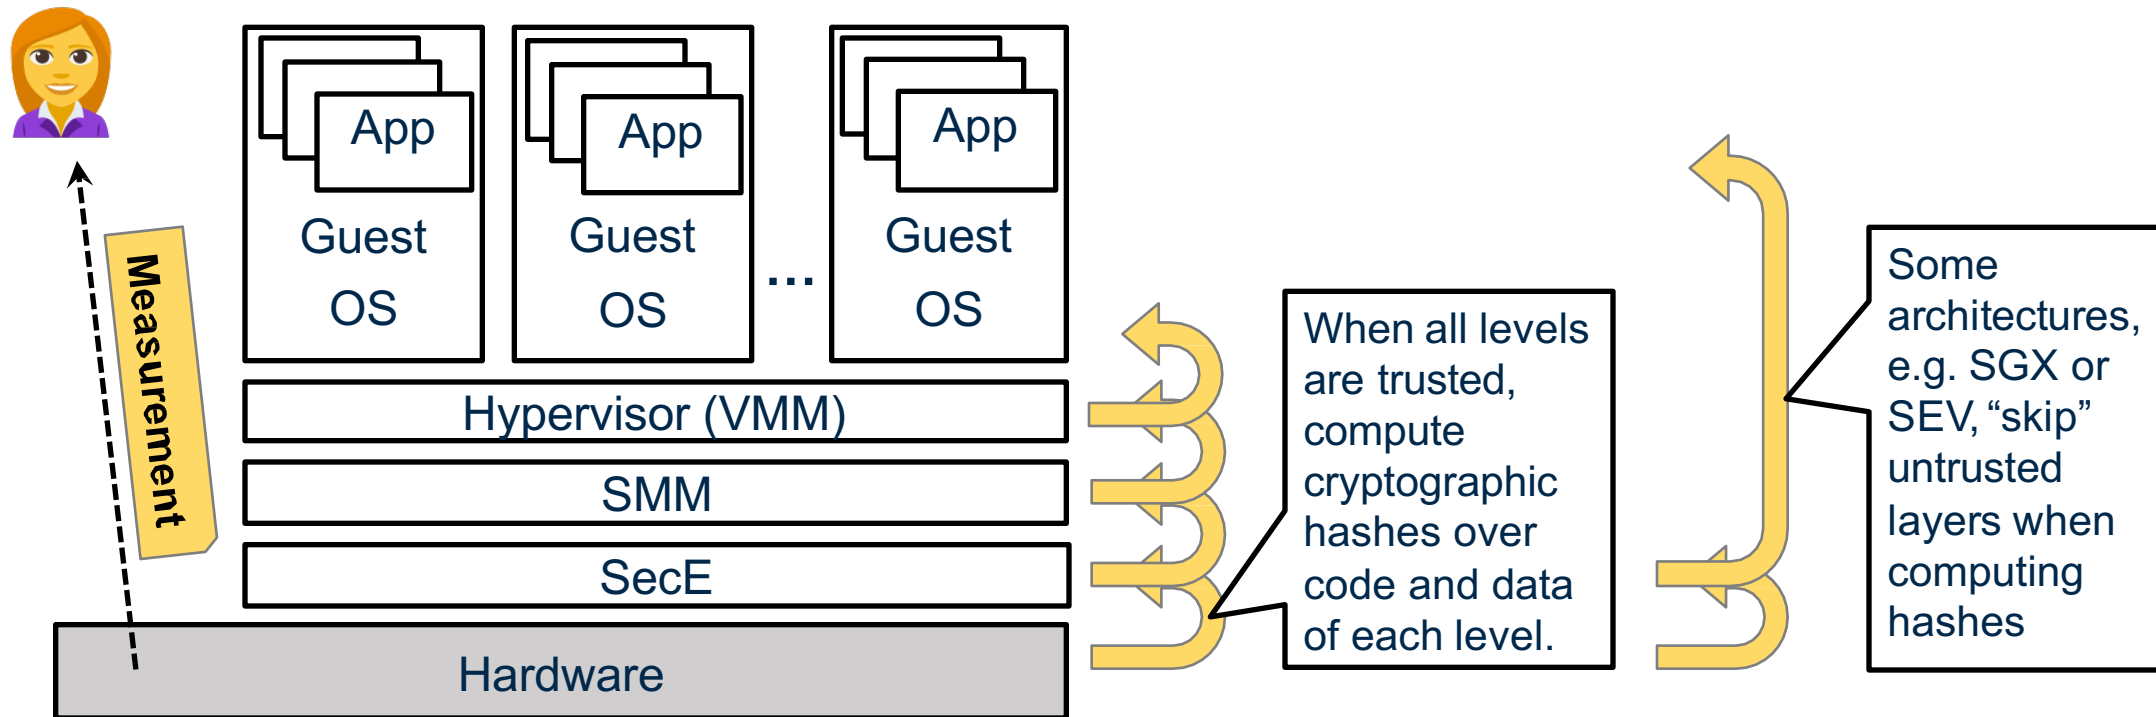
The root of trust is derived from the signing and verification keys.

- › Public key,  $K_{PK}$ , for encrypting data to be sent to the processor
  - › Data handled by the TCB
- › Signature verification key,  $K_{VK}$ , for checking data signed by the processor
  - › TCB can sign user keys
- › Key distribution for PUF-based designs will be different
- › Need infrastructure!



# Software Measurement

With an embedded signing key, the software running in the TEE can be “measured” to attest to external users what code is running on the system.



# Using Software Measurement

## Trusted / Secure / Authenticated Boot:

- › Abort boot when the wrong measurement is obtained
- › Continue booting, but do not decrypt secrets
- › Legitimate software updates will change measurements and may prevent correct boot up

## Remote attestation:

- › Measure and digitally sign measurements that are sent to remove user

## Data sealing (local or remote):

- › Only unseal data if correct measurements are obtained

## TOC-TOU attacks and measurements:

- › Time-of-Check to Time-of-Use (TOC-TOU) attacks leverage the delay between when a measurement is taken and when the component is used
- › Cannot easily use hashes to prevent TOC-TOU attacks

# Need for Continuous Monitoring of Protected Software

Continuous monitoring is a potential solution to TOC-TOU:

- › Constantly measure the system, e.g., performance counters, and look for anomalies
- › Requires knowing the correct and expected behavior of the system
- › It can be used for continuous authentication

Attackers can “hide in the noise” if they change the execution of the software slightly and do not affect performance counters significantly.



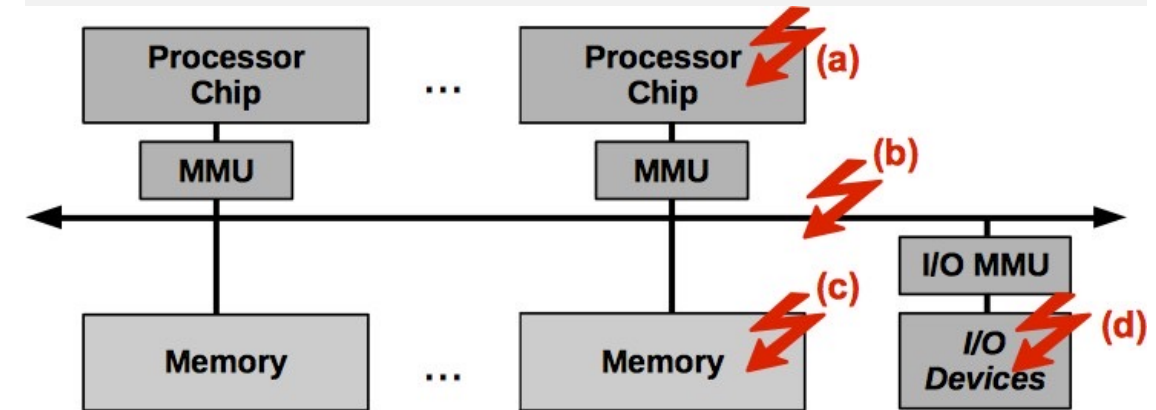
Politecnico  
di Torino

# System Memory Protection

# Sources of Attacks on Memory

Memory is vulnerable to different types of attacks:

- › Untrusted software running on the processor
- › Physical attacks on the memory bus, other devices snooping on the bus, man-in-the-middle attacks with malicious device
- › Physical attacks on the memory (Coldboot, ...)
- › Malicious devices using DMA or other attacks



# Types of Attacks on Memory

Different types of attacks exist (very similar to attacks in network settings):

- Snooping

Passive attack, try to read data contents.

- Spoofing

Active attack, inject new memory commands to try to read or modify data.

- Splicing

Active attack, combine portions of legitimate memory commands into new memory commands (to read or modify data).

- Replay

Active attack, re-send old memory command (to read or modify data).

- Disturbance

Active attack, DoS on memory bus, repeated memory accesses to age circuits, repeated access to make Rowhammer, etc.

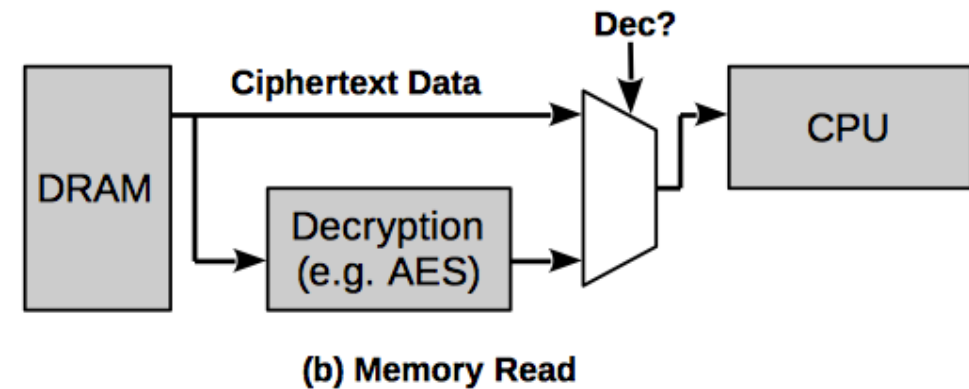
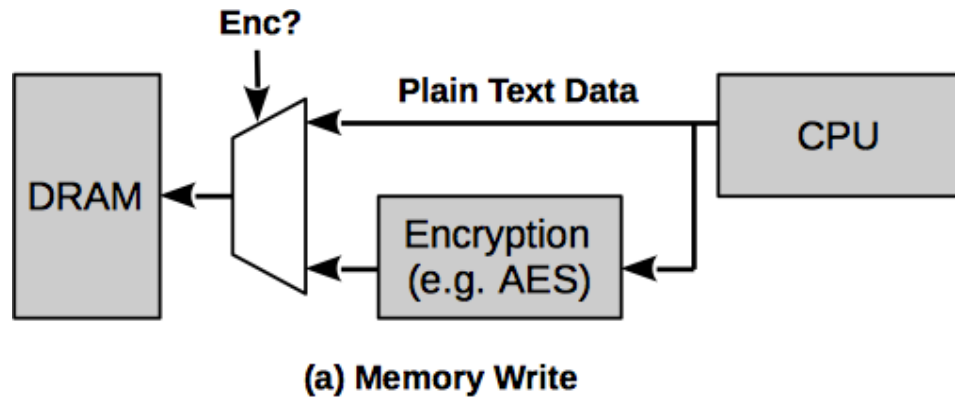
# Confidentiality Protection with Encryption

Contents of the memory can be protected with encryption. Data going out of the CPU is encrypted, and data coming from memory is decrypted before CPU use.

a) Encryption engine (usually AES in CTR mode) encrypts data going out of the processor chip

b) The decryption engine decrypts incoming data

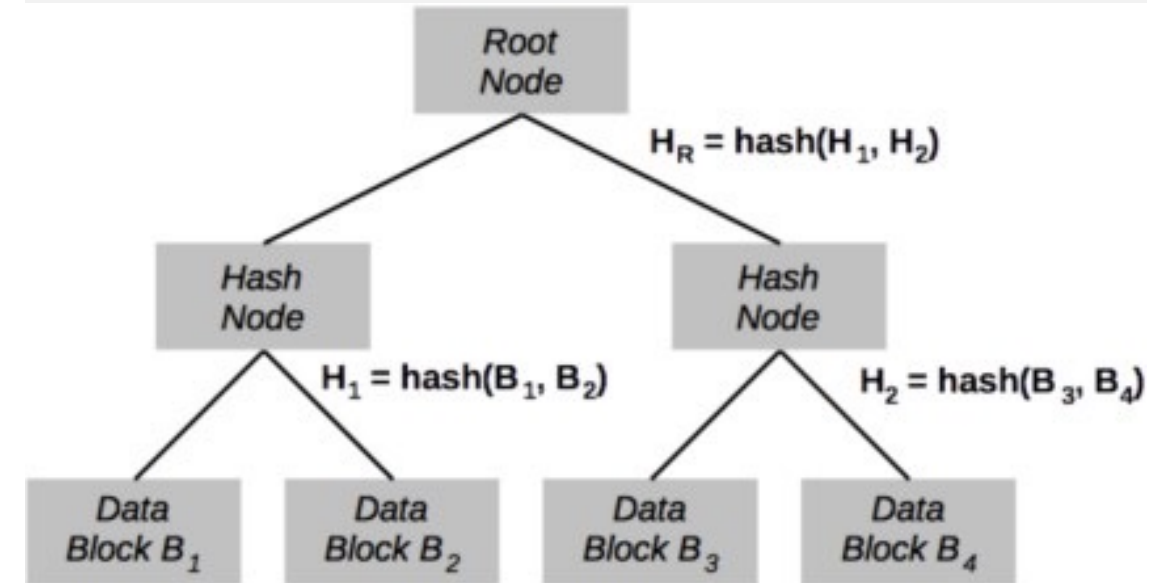
Pre-compute encryption pads, then only need to do XOR; speed depends on how well counters are fetched / predicted.





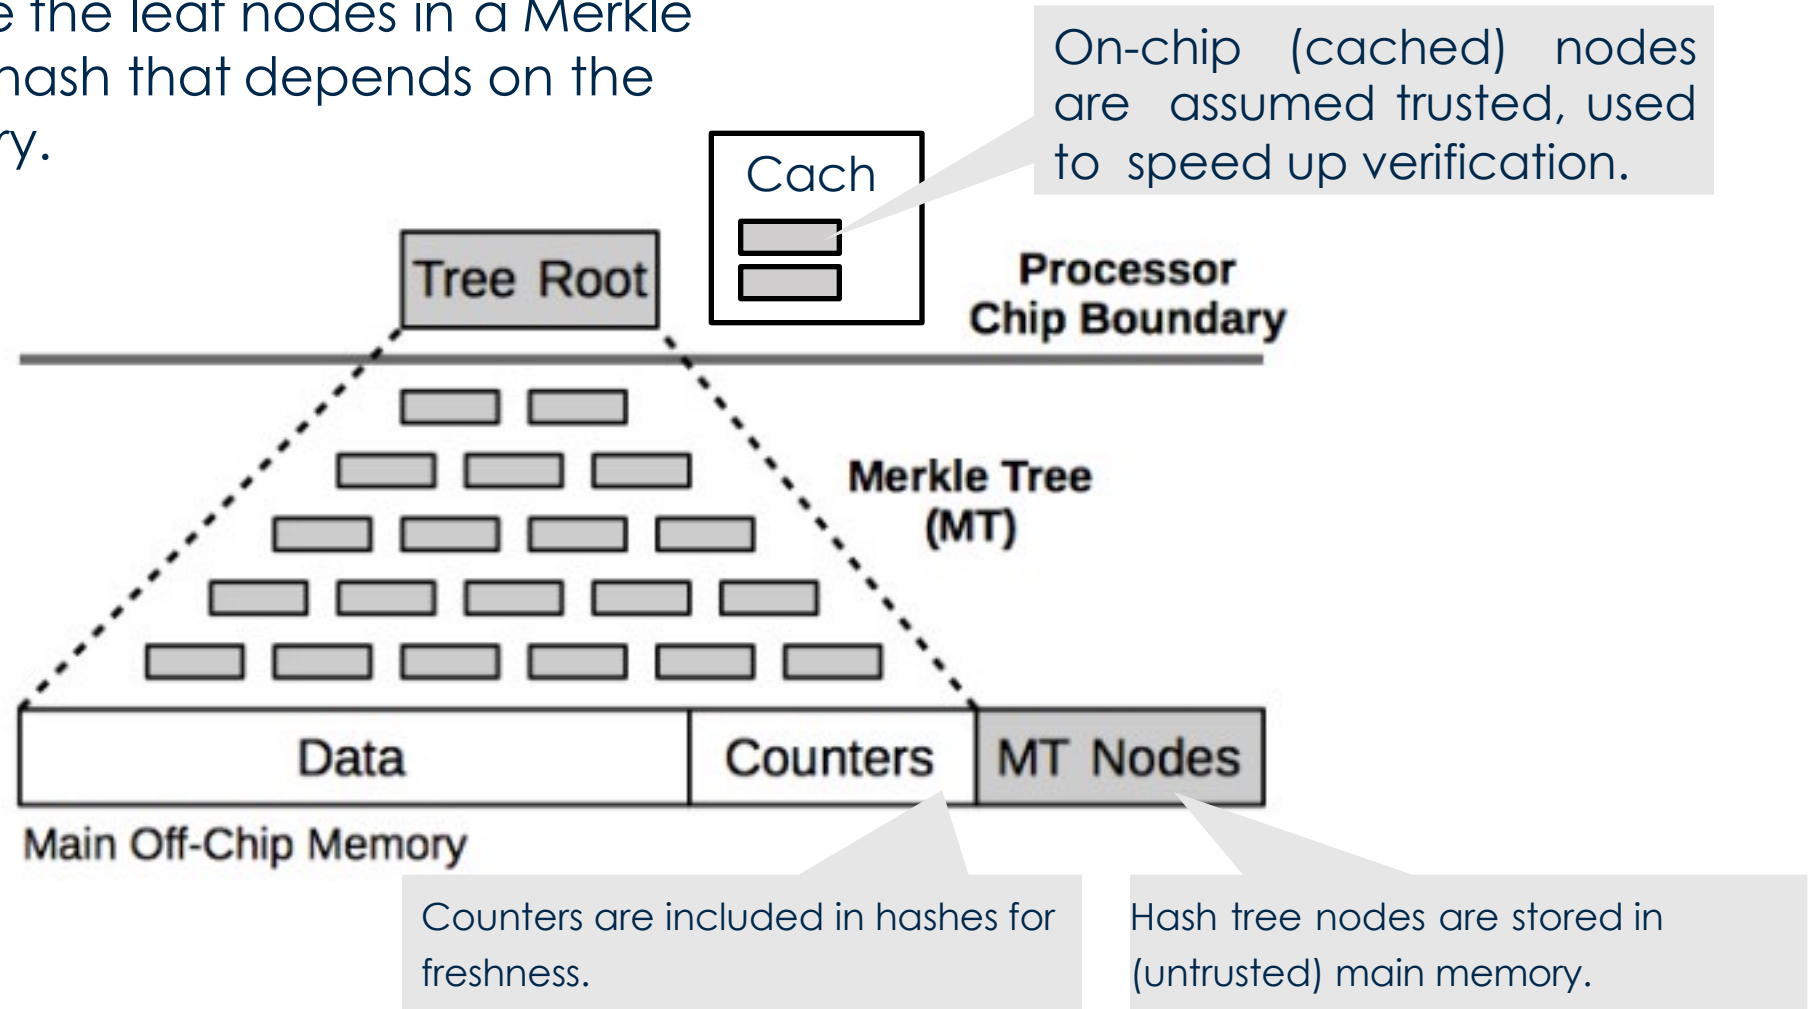
# Hash Tree

A hash tree (also called Merkle Tree) is a logical tree structure, typically a binary tree, where two child nodes are hashed together to create a parent node; the root node is a hash that depends on the value of all the leaf nodes.



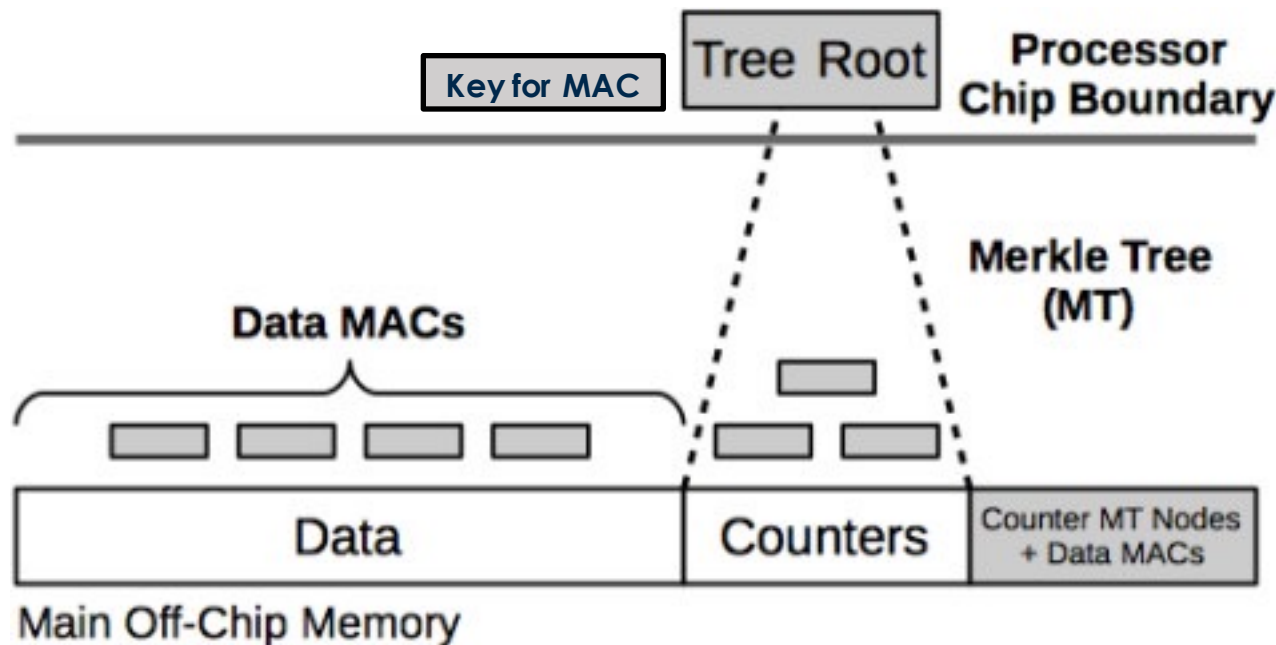
# Integrity Protection with Hash Trees

Memory blocks can be the leaf nodes in a Merkle Tree. The tree root is a hash that depends on the contents of the memory.



# Integrity Protection with Bonsai Hash Trees

Message Authentication Codes (MACs) can be used instead of hashes, and a smaller “Bonsai” tree can be constructed.




# Integrity Protection of Selected Memory Regions


For encryption, the type of encryption does not typically depend on memory configuration

For integrity, the integrity tree needs to consider:

- Protect whole memory
- Protect parts of memory (e.g., per application, VM, etc.)
- Protect external storage (e.g., data swapped to disk)



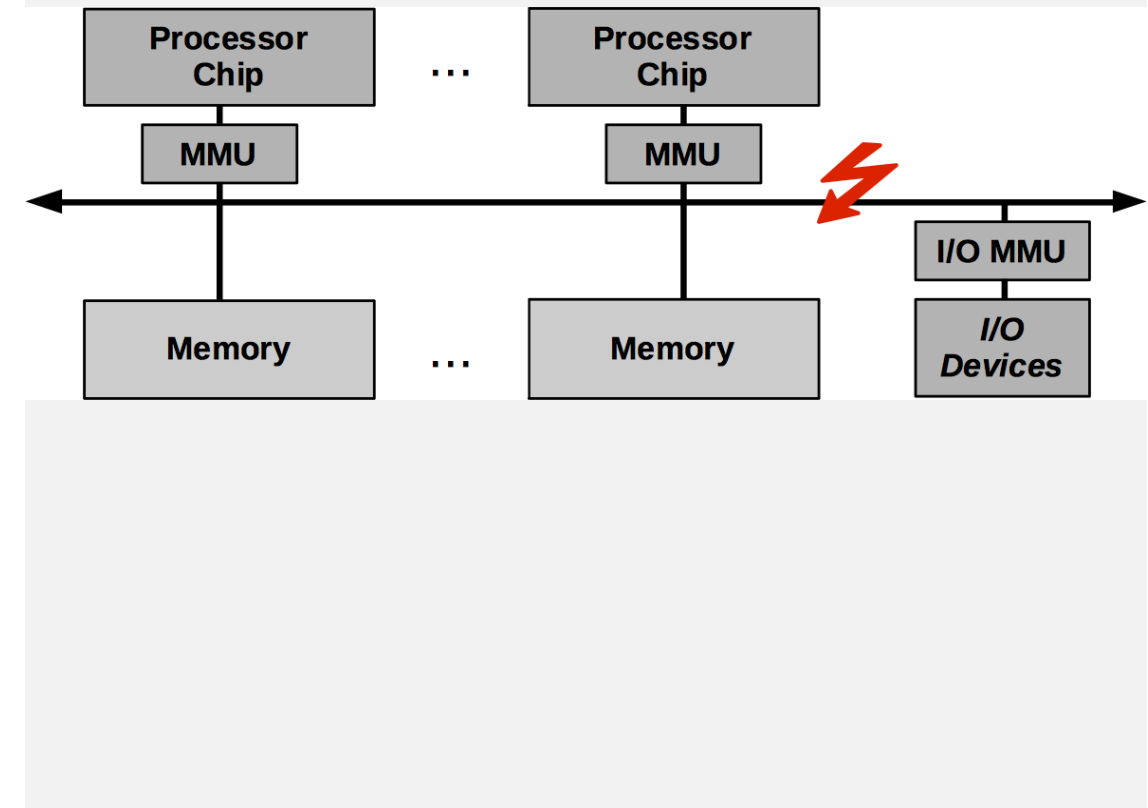
E.g., Bastion's memory integrity tree (Champagne et al., HPCA '10)



# Memory Access Pattern Protection

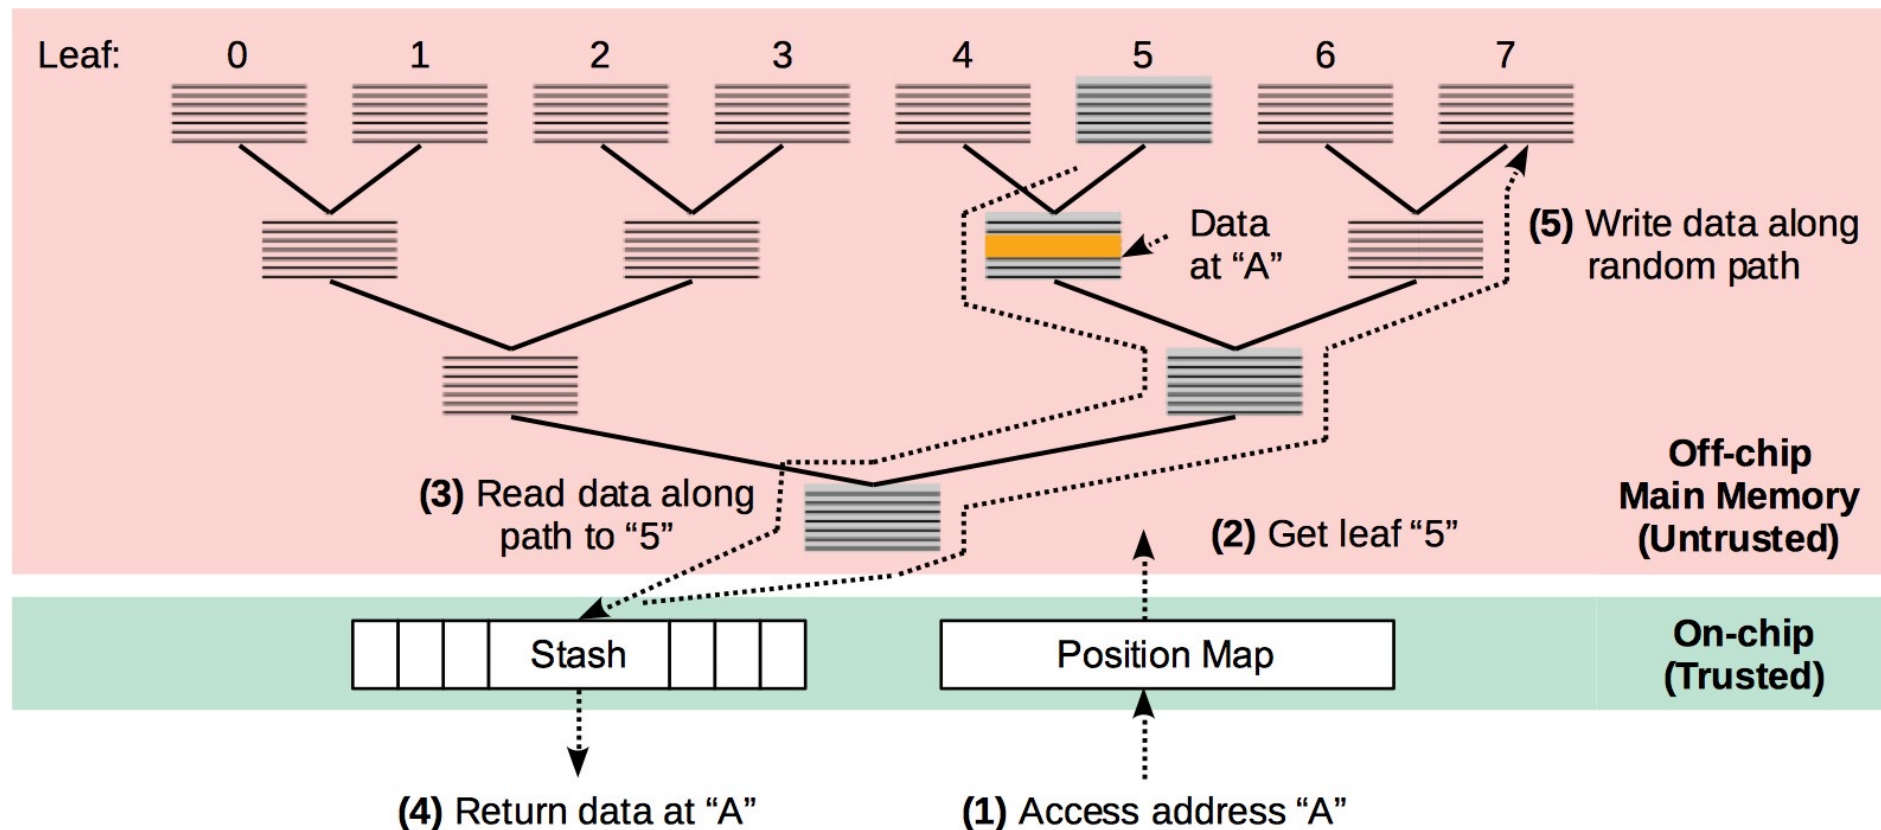
Snooping attacks can target extracting data (protected with encryption) or extracting access patterns to learn what a program is doing.

- › Easier in Symmetric multiprocessing (SMP) due to shared bus
- › Possible in other configurations if there are untrusted components



# Access Patterns Protection


Access patterns (traffic analysis) attacks can be protected using Oblivious RAM, such as Path ORAM. This is on top of encryption and integrity checking.



# Leveraging 2.5D and 3D Integration

With 2.5D and 3D integration, the memory is brought into the same package as the main processor chip. Further, with embedded DRAM (eDRAM), the memory is on the same chip.

- Potentially probing attacks are more difficult
- Still limited memory (eDRAM around 128MB in 2017)



# Security of Non-Volatile Memories and NVRAMs

Non-volatile memories (NVMs) can store data even when there is no power

Non-volatile random-access memory (NVRAM) is a specific type of NVM that is suitable to serve as a computer system's main memory and replace or augment DRAM

- › Many types of NVRAMs:
  - › ReRAM – based on memristors, stores data in the resistance of a dielectric material
  - › FeRAM – uses ferroelectric material instead of a dielectric material
  - › MRAM – uses ferromagnetic materials and stores data in the resistance of a storage cell
  - › PCM – typically uses chalcogenide glass, where different glass phases have different resistances

## Security considerations

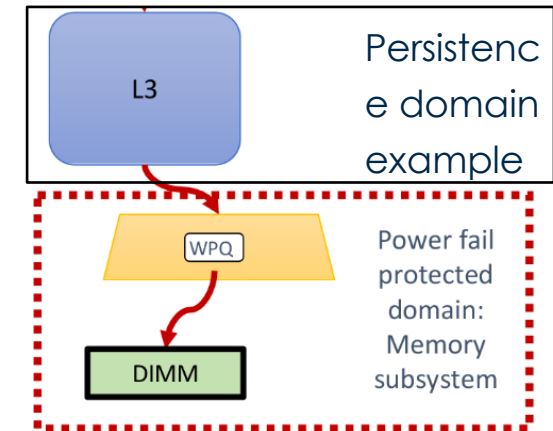
- › Data remanence makes passive attacks easier (e.g., data extraction)
- › Data is maintained after reboot or crash (security state also needs to be correctly restored after reboot or crash)



# Features of Systems using NVRAMs

## Persistence:

- › Data persists across reboots and crashes, possibly with errors
- › Need atomicity for data larger than one memory word (either all data or no data is “persisted”)
  - › E.g., Write Pending Queue (WPQ) – memory controller has non-volatile storage or enough stored charge to write pending data back to the NV-DIMM or NVRAM




## The granularity of persistence:

- › Hide non-volatility from the system: use memory as a DRAM replacement
- › Expose non-volatility to the system: allow users to select which data is non-volatile
- › Linux support through Direct Access (DAX) since about 2014
- › Developed for NV-DIMMs (e.g., battery-backed DRAM, but works for NVRAMs)

Image:  
[https://www.usenix.org/system/files/login/articles/login\\_summer17\\_07\\_rudoff.pdf](https://www.usenix.org/system/files/login/articles/login_summer17_07_rudoff.pdf)

# Integrity Protection of NVRAMs

- For integrity, the integrity tree needs to consider additionally:
  - Atomicity of memory updates for data and related security state (so it is correct after reboot or a crash)
  - Which data in NVRAM is to be persisted (i.e., granularity)



# Encrypted, Hashed, Oblivious Access Memory Assumption

Off-chip memory is untrusted, and the contents are assumed to be protected from snooping, spoofing, splicing, replay, and disturbance attacks:

- › Encryption – snooping and spoofing protection
- › Hashing – spoofing, splicing, replay (counters must be used), and disturbance protection
- › Oblivious Access – snooping protection



Politecnico  
di Torino

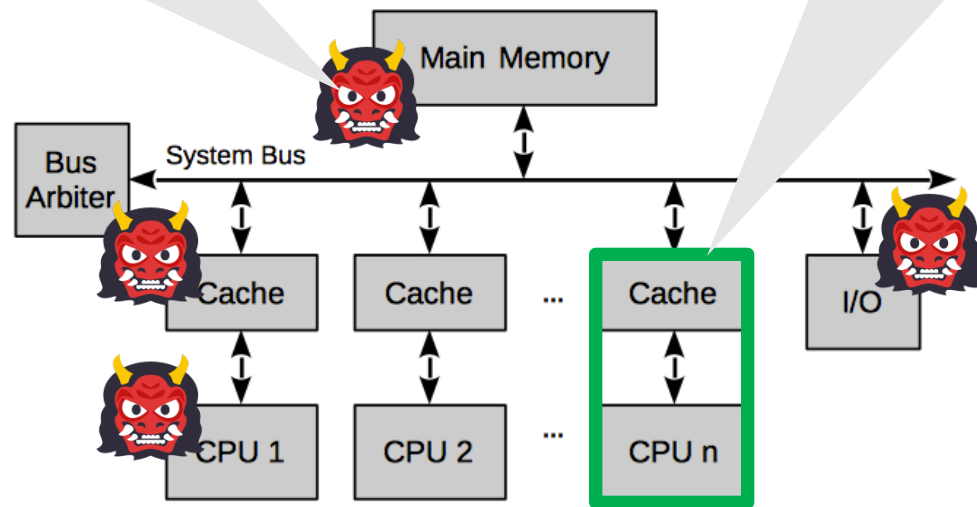
# Multiprocessor and Many-Core protection

# Multiprocessor Architectures

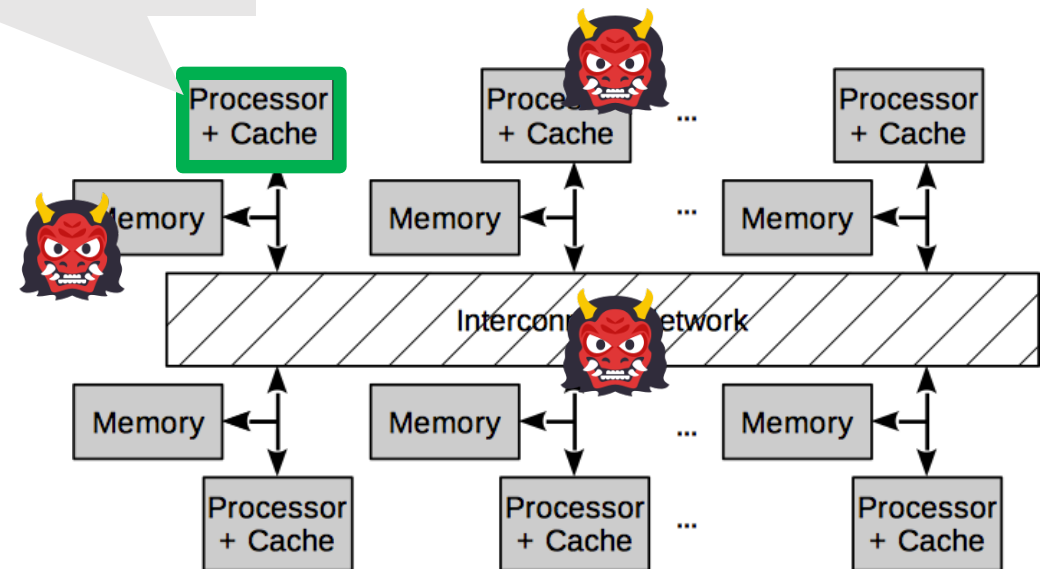
Symmetric Multi-Processing (SMP) and Distributed Share Memory (DSM), also referred to as Non-Uniform Memory Access (NUMA), offer two ways of connecting many CPUs.

Other components on the same system are untrusted

Individual processors are still trusted



SMP



DSM / NUMA

Emoji Image:  
<https://www.emojione.com/emoji/1f479>

# SMP Protections

Encrypt traffic on the bus between processors

- › Each source-destination pair can share a hard-coded key
- › Or use distribute keys using public key infrastructure (within a computer)

Use MACs for the integrity of messages

- › Again, each source-destination pair can share a key

Use Merkle trees for memory protection

- › Can snoop on the shared memory bus to update the tree root node as other processors are doing memory accesses
- › Or per-processor tree

# DSM / NUMA Protections

Encrypt traffic on the bus between processors

- › need a shared key

Use MACs for the integrity of messages

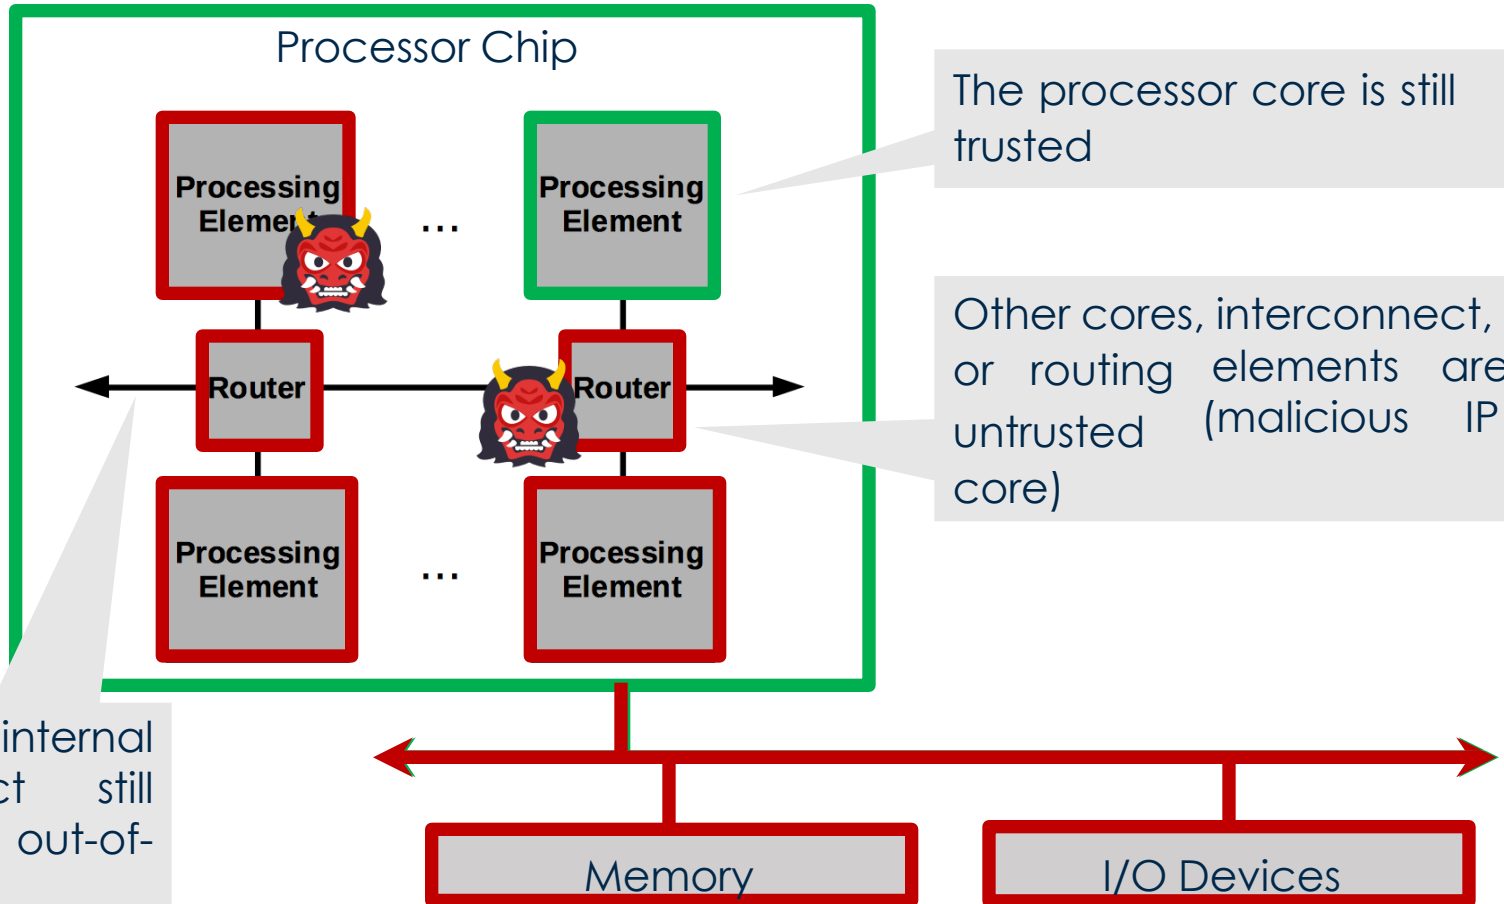
- › each source-destination pair can share a key

Use Merkle trees for memory protection

- › No longer can snoop on the traffic (DSM is point to point usually)

# Many-core Trust Boundary

Trusted processor chip boundary is reduced in most research focusing on many-core security



Emoji Image:  
<https://www.emojione.com/emoji/1f479>



# Architecture and Hardware Security Intersection

With many-core chips, the threats architects worry about start to overlap with hardware security researchers' work

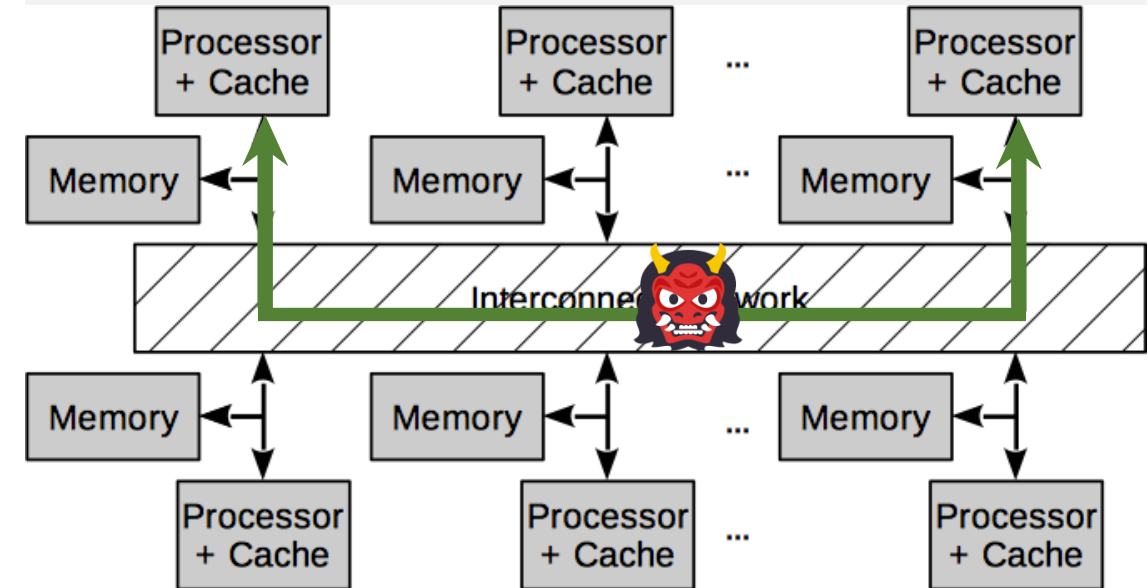
- › Untrusted 3rd party intellectual property (IP) cores
- › Malicious foundry
- › Untrusted supply chain

Architecture solutions (add encryption, add hashing, etc.) complement defenses developed by hardware security experts (split manufacturing, etc.).

# Protected Inter-processor Communication

In addition to the existing assumption about protected memory communication, designs with multiple processors or cores assume the inter-processor communication will be protected:

- › Confidentiality
- › Integrity
- › Communication pattern protection



Emoji Image:

<https://www.emojione.com/emoji/1f479>

# Performance Challenges

Interconnects between processors are very fast:

- › E.g., HyperTransport specifies speeds over 50 GB/s
  - › AES block size is 128 bits
  - › Encryption would need 3 billion (giga) AES block encryptions or decryptions per second
- › Tricks such as counter-mode encryption can help
  - › Only XOR data with a pad
  - › But need to have or predict counters and generate the pads in time



Politecnico  
di Torino

# Secure Design Considerations

# Principles of Secure Processor Architecture Design

Four principles for secure processor architecture design based on existing designs and also on ideas about what ideal design should look like are:

- › Protect Off-chip Communication and Memory
- › Isolate Processor State among TEE Execution and other Software
- › Allow TCB Introspection
- › Authenticate and Continuously Monitor TEE and TCB

Additional design suggestions:

- › Avoid code bloat
- › Minimize TCB
- › Ensure hardware security (Trojan prevention, supply chain issues, etc.)
- › Use formal verification

- Architectural state
- Micro-architectural state
- Due to spatial or temporal sharing of hardware

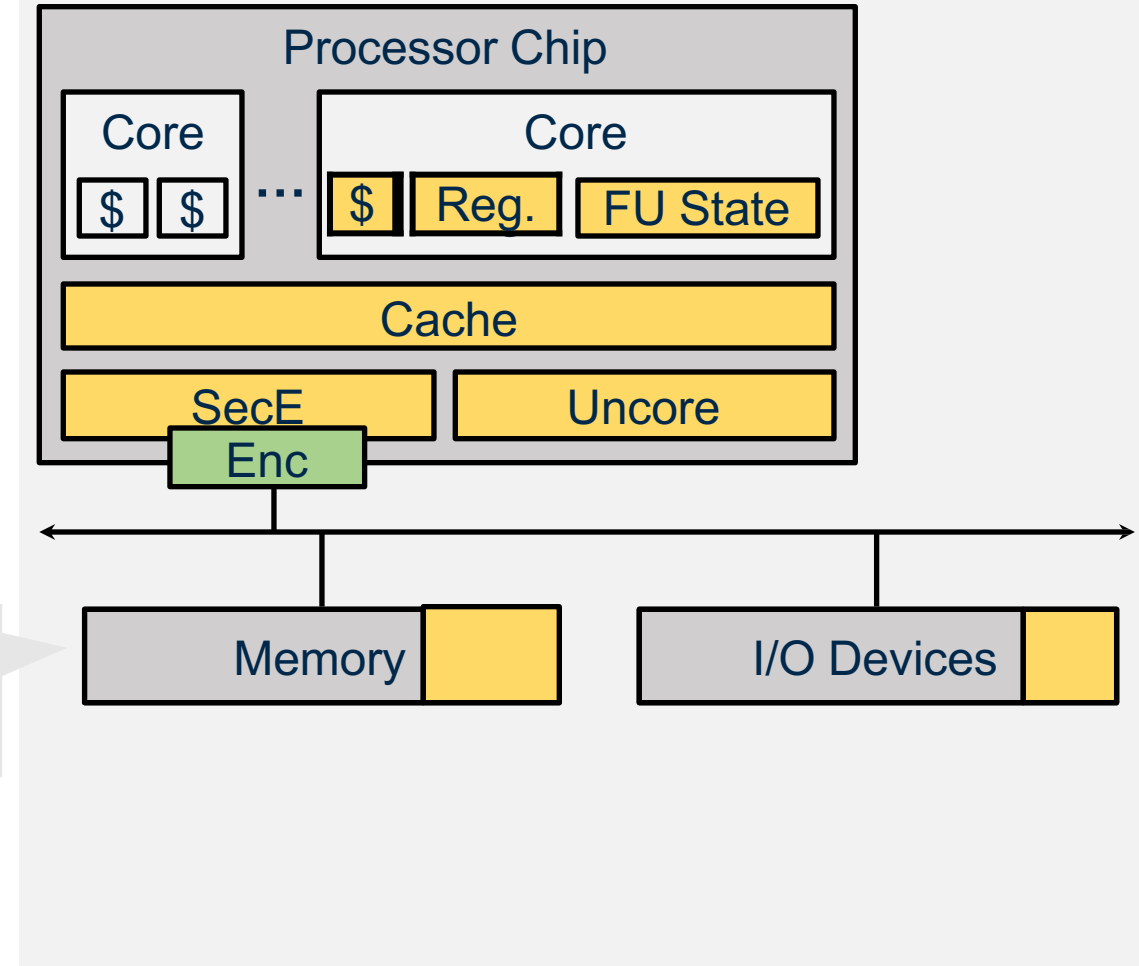
# Protect Off-chip Communication and Memory

Off-chip components and communication are untrusted and need encryption, hashing, and access pattern protection.

Open research challenges:

- › Performance
- › Key distribution

E.g. encryption defends Cold boot style attacks on main memory.

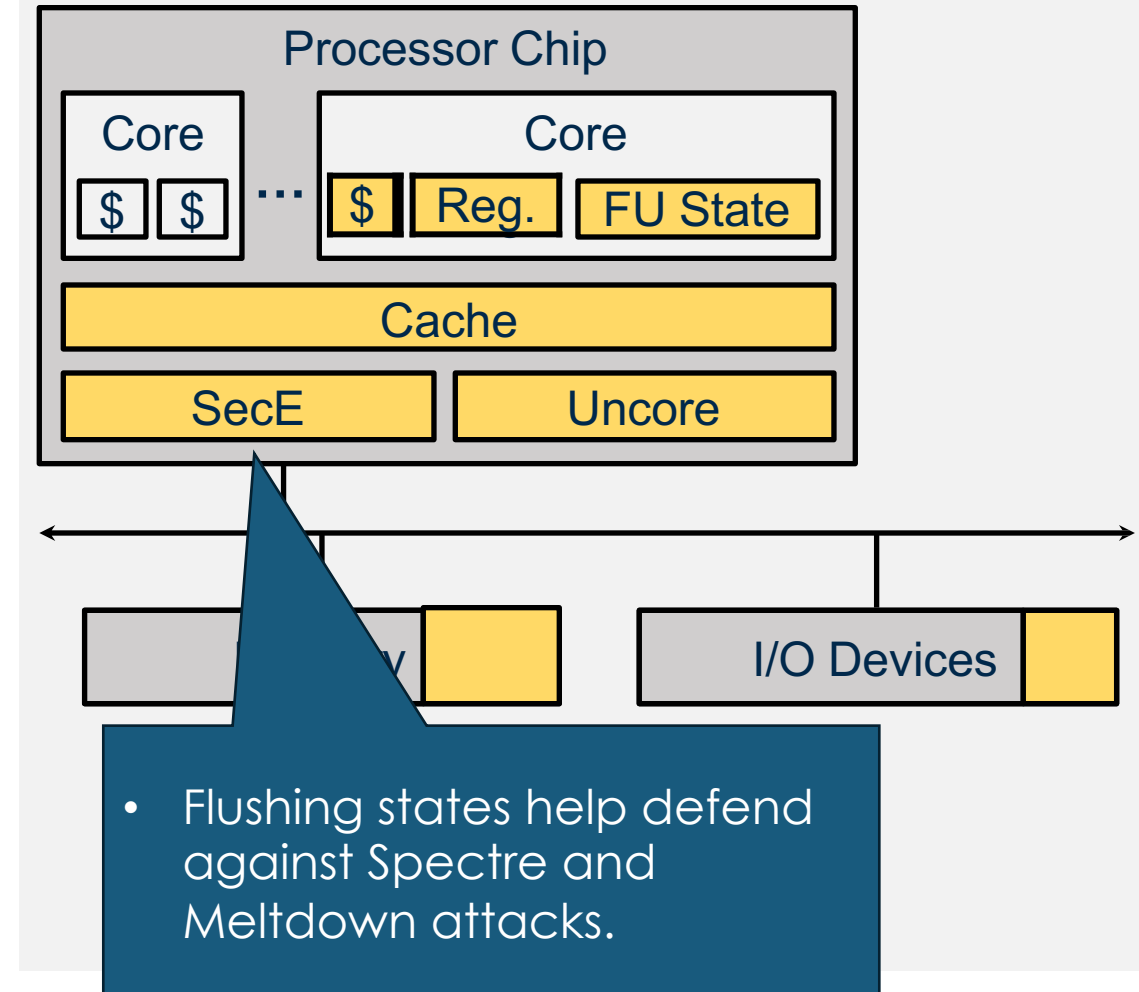


# Isolate Processor State among TEE Execution

When switching between protected software and other software, either protected or not, you must flush the state or save and restore it to prevent one software from influencing another.

Open research challenges:

- › Performance
- › Finding all the states to flush or clean
- › Isolate state during concurrent execution
- › ISA interface to allow state flushing

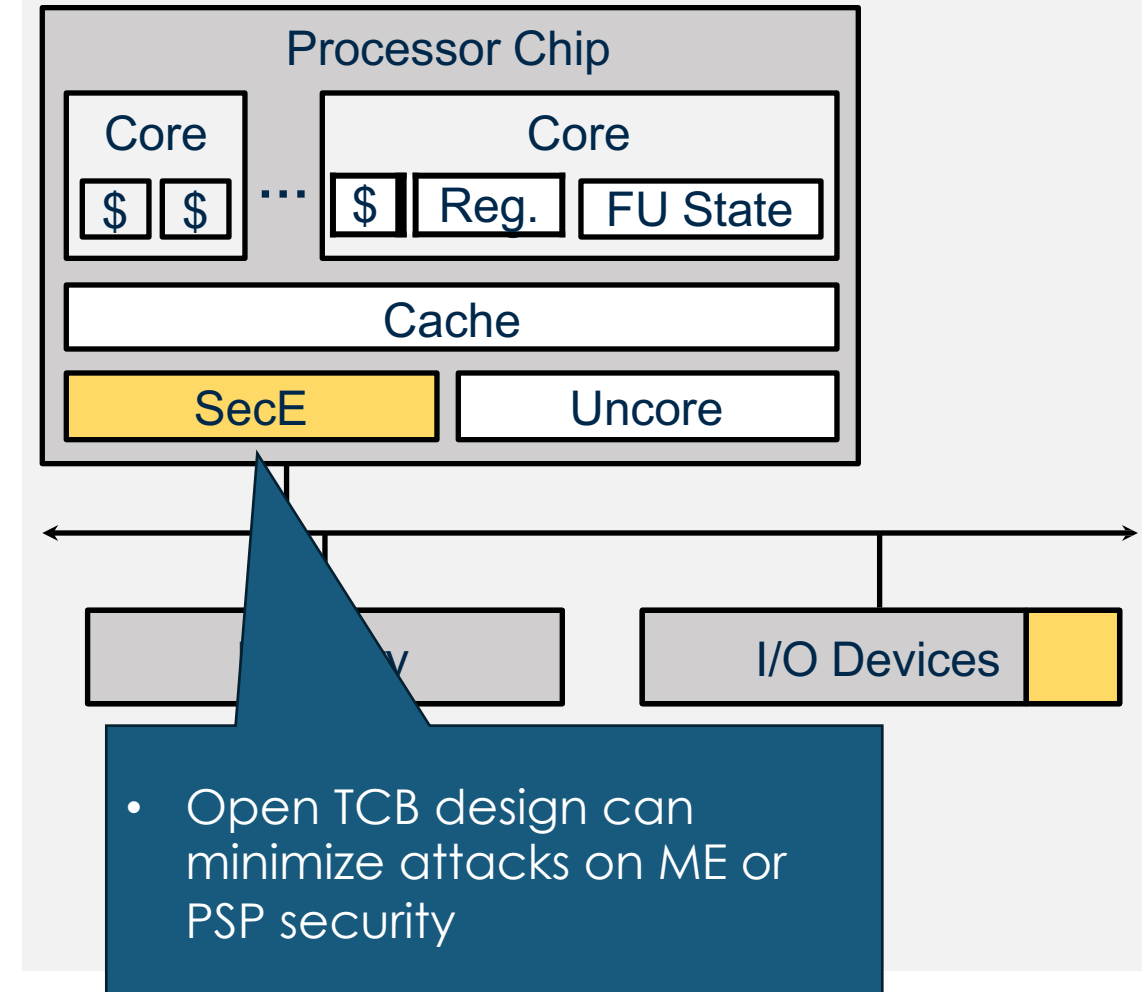


# Allow TCB Introspection

We need to ensure the correct execution of TCB through open access to TCB design, monitoring, fingerprinting, and authentication.

Open research challenges:

- › ISA interface to introspect TCB
- › Area, energy, and performance costs due to extra features for introspection
- › Leaking information about TCB or TEE





# Authenticate and Continuously Monitor TEE and TCB


Monitoring of software running inside TEE, e.g., TSMs or Enclaves, gives assurances about the state of the protected software.

Likewise, monitoring TCB ensures that protections are still in place.

Open research challenges:

- Interface design for monitoring
- Leaking information about TEE

E.g., continuous monitoring of a TEE can help prevent TOC-TOU attacks.



# Pitfalls and Fallacies

## Pitfall: Security by Obscurity

- › E.g., recent attacks on industry processors.

## Fallacy: Hardware Is Immutable

- › Most actually realized architectures use a security processor (e.g., ME or PSP).

## Pitfall: Wrong Threat Model

- › For example, the original SGX did not claim side-channel protection, but researchers attacked it.

## Pitfall: Fixed Threat Model

- › Most designs are one-size-fits-all solutions.

# Pitfalls and Fallacies

## Pitfall: Use of Outdated or Custom Crypto

- › For example, today's devices will be in the field for many years but will not use post-quantum crypto.

## Pitfall: Not Addressing Side Channels

- › Most architectures underestimate side channels.

## Pitfall: Requiring Zero-Overhead Security

- › Performance-, area-, or energy-only focused designs ignore security.

## Pitfall: Code Bloat

- › For example, rather than partitioning a problem, large code pieces run instead of TEEs; TCB gets bigger and bigger, leading to bugs.

# Pitfalls and Fallacies

## Pitfall: Incorrect Abstraction

- › Abstraction (e.g., ISA assumptions) does not match how the device or hardware behaves.

## Pitfall: Focus Only on Speculative Attacks

- › Defending only speculative attacks does not ensure classical attacks are also protected